

Binary Code Analysis: Techniques, Tools, and Applications

Lecture 1: Introduction

Zhiqiang Lin

Department of Computer Science
University of Texas at Dallas

JOSS

Outline

- 1 Background
- 2 Challenges
- 3 Techniques
- 4 Tools
- 5 Applications
- 6 Summary

Acknowledgment: A portion of the slides in this lecture is compiled from presentations by Prof. Tom Reps and also Fish (author of `angr`)

Outline

- 1 Background
- 2 Challenges
- 3 Techniques
- 4 Tools
- 5 Applications
- 6 Summary

GOSS

What is Binary Analysis

The process of (automatically) reasoning/deriving properties about the structure/behavior/syntactics/semantics/anything of your interest of binary programs

```
zlin@zlin-desktop:~/ $ hexdump -C /bin/ls | less
```

```
00000000  7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00 |.ELF.....|
00000010  02 00 3e 00 01 00 00 00 a4 45 40 00 00 00 00 00 |..>.....E@...|
00000020  40 00 00 00 00 00 00 00 70 96 01 00 00 00 00 00 |@.....p.....|
00000030  00 00 00 00 40 00 38 00 09 00 40 00 1c 00 1b 00 |....@.8...@....|
00000040  06 00 00 00 05 00 00 00 40 00 00 00 00 00 00 00 |.....@.....|
00000050  40 00 40 00 00 00 00 00 40 00 40 00 00 00 00 00 |@.@.....@.@....|
00000060  f8 01 00 00 00 00 00 00 f8 01 00 00 00 00 00 00 |.....|
00000070  08 00 00 00 00 00 00 00 03 00 00 00 04 00 00 00 |.....|
00000080  38 02 00 00 00 00 00 00 38 02 40 00 00 00 00 00 |8.....8.@....|
00000090  38 02 40 00 00 00 00 00 1c 00 00 00 00 00 00 00 |8.@.....|
...
```

Why Binary Code?

Access to the source code often is not possible:

- Proprietary software packages. (Volks Wagon's cheating software)
- Stripped executables.
- Proprietary libraries
- Malicious software (exploits), e.g., stuxnet

GOSS

Why Binary Code?

Access to the source code often is not possible:

- Proprietary software packages. (Volks Wagon's cheating software)
- Stripped executables.
- Proprietary libraries
- Malicious software (exploits), e.g., stuxnet

Binary code is the only authoritative version of the program.

- Binary code is everywhere
- Changes occurring in the compile, optimize and link steps can create non-trivial semantic differences from the source and binary.
- What you see is not what you execute (WYSINWYX problem)

Why Binary Code?

- Windows
 - ▶ Login process keeps a user's password in the heap after a successful login
- To minimize data lifetime
 - ▶ clear buffer
 - ▶ call free()

```
memset(buffer, '\\0', len);  
free(buffer);
```

Why Binary Code?

- Windows

- ▶ Login process keeps a user's password in the heap after a successful login

- To minimize data lifetime

- ▶ clear buffer
- ▶ call free()

- But . . .

- ▶ the compiler might optimize away the buffer-clearing code ("useless-code" elimination)

```
memset(buffer, '\\0', len);  
free(buffer);
```



```
free(buffer);
```


Why Binary Code: Backdoor



Linux embedded device: HTTP server for management and video monitoring, with a known backdoor.

Backdoor!!!

- Username: 3sadmin
- Password: 27988303

```
LDR R1, =a3sadmin ; "3sadmin"
MOV R0, R7 ; s1
BL strcmp
CMP R0, #0
LDR R1, =a27988303 ; "27988303"
MOV R0, R4 ; s1
```

Heffner, Craig. "Finding and Reversing Backdoors in Consumer Firmware." EELive! (2014).

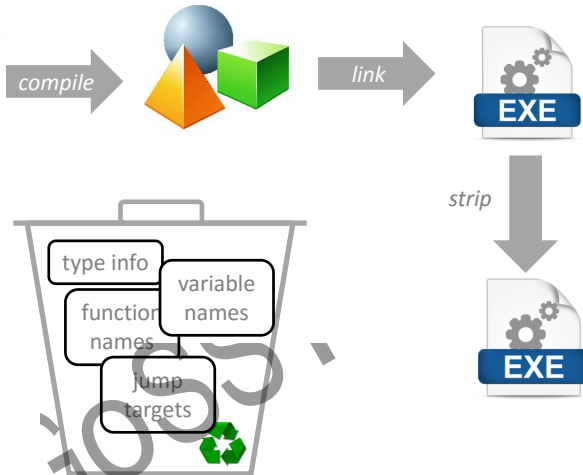
WYSINWYX

What You See Is Not What You eXecute

```

1 // compile
2 // Data is source IR2 output, meaning a IR2output containing the result.
3 // It is meant to be used as an IR2 input to other IR2 input outputs.
4 // compile
5 // compile
6 // compile
7 // compile
8 // compile
9 // compile
10 // compile
11 // compile
12 // compile
13 // compile
14 // compile
15 // compile
16 // compile
17 // compile
18 // compile
19 // compile
20 // compile
21 // compile
22 // compile
23 // compile
24 // compile
25 // compile
26 // compile
27 // compile
28 // compile
29 // compile
30 // compile
31 // compile
32 // compile
33 // compile
34 // compile
35 // compile
36 // compile
37 // compile
38 // compile
39 // compile
40 // compile
41 // compile
42 // compile
43 // compile
44 // compile
45 // compile
46 // compile
47 // compile
48 // compile
49 // compile
50 // compile
51 // compile
52 // compile
53 // compile
54 // compile
55 // compile
56 // compile
57 // compile
58 // compile
59 // compile
60 // compile
61 // compile
62 // compile
63 // compile
64 // compile
65 // compile
66 // compile
67 // compile
68 // compile
69 // compile
70 // compile
71 // compile
72 // compile
73 // compile
74 // compile
75 // compile
76 // compile
77 // compile
78 // compile
79 // compile
80 // compile
81 // compile
82 // compile
83 // compile
84 // compile
85 // compile
86 // compile
87 // compile
88 // compile
89 // compile
90 // compile
91 // compile
92 // compile
93 // compile
94 // compile
95 // compile
96 // compile
97 // compile
98 // compile
99 // compile
100 // compile

```



An Example of WYSINWYX

```
int callee(int a, int b) {  
    int local;  
    if (local == 5) return 1;  
    else return 2;  
}
```

```
int main() {  
    int c = 5;  
    int d = 7;
```

```
    int v = callee(c,d);  
    // What is the value of v here?  
    return 0;  
}
```

An Example of WYSINWYX

```
int callee(int a, int b) {  
    int local;  
    if (local == 5) return 1;  
    else return 2;  
}
```

Answer: 1
(for the Microsoft compiler)

```
int main() {  
    int c = 5;  
    int d = 7;
```

```
    int v = callee(c,d);  
    // What is the value of v here?  
    return 0;  
}
```

Tutorial on x86 (Intel Syntax)

```
p = q;
```

```
p = *q;
```

```
*p = q;
```

```
p = &a[2];
```

GOSS

Tutorial on x86 (Intel Syntax)

```
ecx = edx;
```

```
ecx = *edx;
```

```
*ecx = edx;
```

```
ecx = &a[2];
```

GOSS

Tutorial on x86 (Intel Syntax)

```
mov    ecx, edx
```

```
ecx = edx;
```

```
ecx = *edx;
```

```
*ecx = edx;
```

```
ecx = &a[2];
```

GOSS

Tutorial on x86 (Intel Syntax)

```
mov
```

```
ecx, edx
```

```
ecx = edx;
```

```
ecx = *edx;
```

```
*ecx = edx;
```

```
ecx = &a[2];
```

GOSS

Tutorial on x86 (Intel Syntax)

```
mov
```

```
ecx, edx
```

```
mov
```

```
ecx, [edx]
```

```
ecx = edx;
```

```
ecx = *edx;
```

```
*ecx = edx;
```

```
ecx = &a[2];
```

GOSS

Tutorial on x86 (Intel Syntax)

```
mov    ecx, edx
mov    ecx, [edx]
mov    [ecx], edx
```

```
ecx = edx;
ecx = *edx;
*ecx = edx;
ecx = &a[2];
```

JOSS

Tutorial on x86 (Intel Syntax)

```
mov    ecx, edx
mov    ecx, [edx]
mov    [ecx], edx
lea    ecx, [esp+8]
```

```
ecx = edx;
ecx = *edx;
*ecx = edx;
ecx = &a[2];
```

JOSS

Tutorial on x86 (Intel Syntax)

```
mov    ecx, edx
mov    ecx, [edx]
mov    [ecx], edx
lea    ecx, [esp+8]
```

```
ecx = edx;
ecx = *edx;
*ecx = edx;
ecx = &a[2];
```

Stack pointer: **esp**

Frame pointer: **ebp**

An Example of WYSINWYX

```
int callee(int a, int b) {  
    int local;  
    if (local == 5) return 1;  
    else return 2;  
}
```

Answer: 1
(for the Microsoft compiler)

```
int main() {  
    int c = 5;  
    int d = 7;  
  
    int v = callee(c,d);  
    // What is the value of v here?  
    return 0;  
}
```

An Example of WYSINWYX

```
int callee(int a, int b,
int local;
if (local == 5) return 1;
else return 2;
}
```

Standard prolog

```
push    ebp
mov     ebp, esp
sub     esp, 4
```

Prolog for 1 local

```
push    ebp
mov     ebp, esp
push    ecx
```

Answer: 1
(for the Microsoft compiler)

```
int main() {
    int c = 5;
    int d = 7;

    int v = callee(c,d);
    // What is the value of v here?
    return 0;
}
```

An Example of WYSINWYX

Standard prolog

push ebp

mov ebp, esp


sub esp, 4

GOSS

An Example of WYSINWYX

Standard prolog

```
push    ebp
mov     ebp, esp
sub     esp, 4
```



GOSS

An Example of WYSINWYX

Standard prolog

```
push    ebp
```


```
mov     ebp, esp
```

```
sub     esp, 4
```

ecx:

5

ebp →



esp →

goss

An Example of WYSINWYX

Standard prolog


```
push    ebp
```

```
mov     ebp, esp
```

```
sub     esp, 4
```

ecx:

5



An Example of WYSINWYX

Standard prolog


push ebp

mov ebp, esp

sub esp, 4

ecx:

5



An Example of WYSINWYX

Standard prolog


push ebp

mov ebp, esp

sub esp, 4

ecx:

5



An Example of WYSINWYX

Standard prolog


push ebp

mov ebp, esp

sub esp, 4

ecx:

5



An Example of WYSINWYX

Standard prolog


push ebp

mov ebp, esp

sub esp, 4

ecx:

5



An Example of WYSINWYX

Standard prolog


```
push    ebp
mov     ebp, esp
sub     esp, 4
```

Prolog for 1 local

```
push    ebp
mov     ebp, esp
push    ecx
```


ecx:

5



ecx:

5



An Example of WYSINWYX


Standard prolog

```
push  ebp
mov   ebp, esp
sub   esp, 4
```


Prolog for 1 local

```
push  ebp
mov   ebp, esp
push  ecx
```

ecx: 5



ecx: 5



An Example of WYSINWYX

```
int callee(int a, int b) {
    int local;
    if (local == 5) return 1;
    else return 2;
}
```

Standard prolog

```
push    ebp
mov     ebp, esp
sub     esp, 4
```

Prolog for 1 local

```
push    ebp
mov     ebp, esp
push    ecx
```

Answer: 1
(for the Microsoft compiler)

```
int main() {
    int c = 5;
    int d = 7;

    int v = callee(c,d);
    // What is the value of v here?
    return 0;
}
```

An Example of WYSINWYX

```
int callee(int a, int b) {
    int local;
    if (local == 5) return 1;
    else return 2;
}
```

Standard prolog

```
push    ebp
mov     ebp, esp
sub     esp, 4
```

Prolog for 1 local

```
push    ebp
mov     ebp, esp
push    ecx
```

Answer: 1
(for the Microsoft compiler)

```
int main() {
    int c = 5;
    int d = 7;

    int v = callee(c,d);
    // What is the value of v here?
    return 0;
}
```

```
mov     [ebp+var_8], 5
mov     [ebp+var_C], 7
mov     eax, [ebp+var_C]
push    eax
mov     ecx, [ebp+var_8]
push    ecx
call    _callee
...
```

An Example of WYSINWYX

```
int callee(int a, int b) {
    int local;
    if (local == 5) return 1;
    else return 2;
}
```

Standard prolog

```
push    ebp
mov     ebp, esp
sub     esp, 4
```

Prolog for 1 local

```
push    ebp
mov     ebp, esp
push    ecx
```

Answer: 1
(for the Microsoft compiler)

```
int main() {
    int c = 5;
    int d = 7;

    int v = callee(c,d);
    // What is the value of v here?
    return 0;
}
```

```
mov     [ebp+var_8], 5
mov     [ebp+var_C], 7
mov     eax, [ebp+var_C]
push    eax
mov     ecx, [ebp+var_8]
push    ecx
call    _callee
...
```

Outline

- 1 Background
- 2 Challenges**
- 3 Techniques
- 4 Tools
- 5 Applications
- 6 Summary

GOSS

What to Reason About in Binary Code?

The process of (automatically) reasoning/deriving properties about the structure/behavior/syntax/semantics/anything of your interest of binary programs

- 1 What are the program's variables and their types?
- 2 What are the program's parameters and their types?
- 3 Where could this indirect jump go?
- 4 What function could be called at this indirect call site?
- 5 What could this dereference operation access/affect?
- 6 What kind of object is allocated at this allocation site?
- 7 What could the value held in V eventually affect?
- 8 What could have affected the value of V ?
- 9 What are the statements (at instruction level) that contribute to the execution of i ?
- 10 ...

Challenges: Abstraction Recovery

The first step in any binary code analysis is to **reconstruct the abstractions distilled after compilation**, such as recognizing the instructions, operand, opcode, variables, basic blocks, and control flows

GOSS

Challenges: Abstraction Recovery

The first step in any binary code analysis is to **reconstruct the abstractions distilled after compilation**, such as recognizing the instructions, operand, opcode, variables, basic blocks, and control flows

Challenges

- Code/Data distinction
- Variable x86 instruction size
- Indirect Branches
- Functions without explicit CALL
- Position independent code (PIC)
- ...

It will be easier to recover these abstractions by using **dynamic analysis**, but will be much more challenge in **static analysis**.

Challenges: Path Coverage, and Path Explosion

For both static analysis and dynamic analysis, how to model the program execution path (too conservative, or too simple), and how to trigger the program path (especially for dynamic analysis) is another key challenge

GOSS

Challenges: Path Coverage, and Path Explosion

For both static analysis and dynamic analysis, how to model the program execution path (too conservative, or too simple), and how to trigger the program path (especially for dynamic analysis) is another key challenge

Static analysis

- Too conservative
- Too many paths
- Impossible path

Dynamic analysis

- A single path
- Cover more path
- Test case generation

A Surprise:

Analyzing Executables can be **Less Complicated** than Analyzing Source

Many source-level issues gone

- 1 Use of multiple source languages
- 2 In-lined assembly code
- 3 Avoid building problems
- 4 Analyze the actual libraries
- 5 ...

Many people inspecting binaries in the whole life

- 1 IDA Pro Users
- 2 Anti-malware companies
- 3 Computer Emergency Response Teams
- 4 Malware writers
- 5 ...

A Surprise:

Executables can be a **Better Platform for Finding Security Vulnerabilities**

- Many exploits utilize platform-specific quirks
 - ▶ non-obvious and unexpected
 - ▶ compiler artifacts (choices made by the compiler)
 - ★ Memory layout
 - padding between fields of a struct
 - which variables are adjacent
 - ★ register usage
 - ★ execution order
 - ★ optimizations performed
 - ★ compiler bugs

Outline

- 1 Background
- 2 Challenges
- 3 Techniques**
- 4 Tools
- 5 Applications
- 6 Summary

GOSS

Basic Techniques

1 Data Flow Analysis

- ▶ Data dependency
- ▶ Taint analysis
- ▶ Point-to analysis

2 Control Flow Analysis

- ▶ Control flow graph
- ▶ Call graph
- ▶ Control dependency

3 Program Slicing

4 Symbolic Execution

Outline

- 1 Background
- 2 Challenges
- 3 Techniques
- 4 Tools**
- 5 Applications
- 6 Summary

GOSS

Static Analysis, Dynamic Analysis, Symbolic Execution

BAP BAT CodeReason
radare2
vivisect Hex-Ray IDA rdis Valgrind
amoco fuzzgrind gdb SemTrax
angr
JARVIS BitBlaze
BARF
klee/s2e Jakstab
insight
PIN QEMU Bindead
Triton
PySysEmu TEMU PEMU
miasm CodeSurfer
paimei

Common Tools

- 1 Static analysis
 - ▶ IDA Pro, BinNav
 - ▶ BAP
- 2 Dynamic analysis
 - ▶ PIN
 - ▶ QEMU
 - ▶ PEMU
- 3 Symbolic execution
 - ▶ FuzzBall, Fuzzgrind
 - ▶ S2E
 - ▶ Angr

Outline

- 1 Background
- 2 Challenges
- 3 Techniques
- 4 Tools
- 5 Applications**
- 6 Summary

GOSS

Applications of Binary Analysis in Security

Use Cases

- 1 Reverse engineering (knowing the secrets)
- 2 Vulnerability discovery/fuzzing
- 3 Exploit generation
- 4 Software verification
- 5 Program testing
- 6 ...

Applications: Vulnerability Discovery



Vulnerability Discovery

How do I trigger path X or condition Y?

Basic Approaches

1 Static Analysis

- ▶ “You can’t ” / “You might be able to”
- ▶ Based on various static techniques.

2 Dynamic Analysis

- ▶ Input A? Input B? Input C? ...
- ▶ Based on concrete inputs to application

3 Symbolic Execution

- ▶ Interpret the application.
- ▶ Track “constraints” on variables.
- ▶ When the required condition is triggered, “concretize” to obtain a possible input

Outline

- 1 Background
- 2 Challenges
- 3 Techniques
- 4 Tools
- 5 Applications
- 6 Summary**

JOSS

Binary Analysis

- Binary code is everywhere, and it is the final representation of programs
- Binary analysis is challenging
- It is extremely useful to perform binary code analysis (vulnerability excavation, backdoor identification) in security
- Basic binary analysis approaches: static/dynamic analysis, symbolic execution
- There are many public available binary analysis tools

Lecture 2: Dynamic Binary Analysis

Zhiqiang Lin

Department of Computer Science
University of Texas at Dallas

JOSS

Outline

1 Basic Concepts

2 QEMU

3 PIN

4 Summary

GOSS

What Is Instrumentation

A technique that inserts extra code into a program to collect information of your interest. Such technique has been widely used in practice in both program debugging and security analysis.

GOSS

What Is Instrumentation

A technique that inserts extra code into a program to collect information of your interest. Such technique has been widely used in practice in both program debugging and security analysis.

```
Max = 0;
for (p = head; p; p = p->next)
{
    printf("In loop\n");
    if (p->value > max)
    {
        printf("True branch\n");
        max = p->value;
    }
}
```

What Is Instrumentation

A technique that inserts extra code into a program to collect information of your interest. Such technique has been widely used in practice in both program debugging and security analysis.

```
Max = 0;
for (p = head; p; p = p->next)
{
    count[0]++;
    if (p->value > max)
    {
        count[1]++;
        max = p->value;
    }
}
```

What Is (Dynamic) Binary Instrumentation

A technique that inserts extra code into the binary code of a program to collect (runtime) information of your interest.

GOSS

What Is (Dynamic) Binary Instrumentation

A technique that inserts extra code into the binary code of a program to collect (runtime) information of your interest.

```
icount++  
sub        $0xff, %edx  
icount++  
cmp        %esi, %edx  
icount++  
jle        <L1>  
icount++  
mov        $0x1, %edi  
icount++  
add        $0x10, %eax
```

What Can Instrumentation Do?

- Profiler for compiler optimization:
 - ▶ Basic-block count
 - ▶ Value profile
- Micro architectural study:
 - ▶ Instrument branches to simulate branch predictors
 - ▶ Generate traces
- Bug checking/Vulnerability identification/Exploit generation:
 - ▶ Find references to uninitialized, unallocated address
 - ▶ Inspect argument at particular function call
 - ▶ Inspect function pointers and return addresses
- Software tools that use dynamic binary instrumentation:
 - ▶ Valgrind, Pin, QEMU, DynInst, ...

Instrumentation approaches: source vs. binary

- Source instrumentation:
 - ▶ Instrument source programs
- Binary instrumentation:
 - ▶ Instrument executables directly
- Advantages for binary instrumentation
 - ▶ Language independent
 - ▶ Machine-level view
 - ▶ Instrument legacy/proprietary software

JOSS

Binary Instrumentation Is Dominant

- Libraries are a big pain for source code level instrumentation
 - ▶ Proprietary libraries: communication (MPI, PVM), linear algebra (NGA), database query (SQL libraries).
- Easily handle multi-lingual programs
 - ▶ Source code level instrumentation is heavily language dependent.
 - ★ More complicated semantics
- Turning off compiler optimizations can maintain an almost perfect mapping from instructions to source code lines
- Worms and viruses are rarely provided with source code
- ...

Instrumentation approaches: static vs. dynamic

- When to instrument
 - ▶ Instrument statically - before runtime
 - ▶ Instrument dynamically - during runtime
- Advantages for dynamic instrumentation
 - ▶ No need to recompile or relink
 - ▶ Discover code at runtime
 - ▶ Handle dynamically-generated code
 - ▶ Attach to running processes

JOSS

How is Instrumentation used in Program Analysis?

- Code coverage
- Call-graph generation
- Memory-leak detection
- Vulnerability identification
- Instruction profiling
- Data dependence profiling
- Thread analysis
 - ▶ Thread profiling
 - ▶ Race detection

GOSS

Outline

1 Basic Concepts

2 QEMU

3 PIN

4 Summary

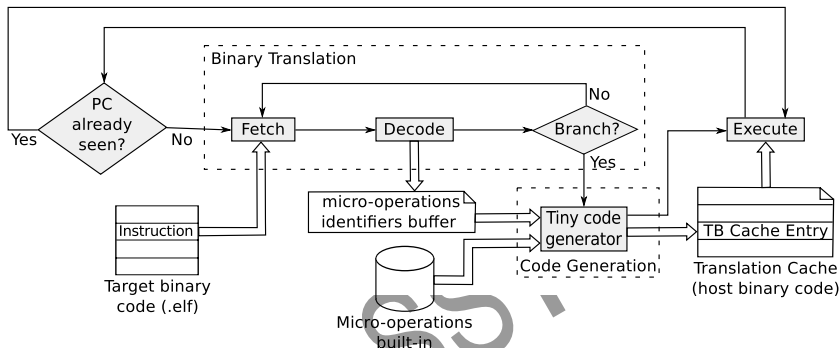
GOSS

QEMU

- QEMU is a generic and open source **machine emulator** and **virtualizer**.
- As a machine **emulator**, QEMU can run OSES and programs made for one machine (e.g. an ARM board) on a different machine (e.g. your own PC). By using dynamic translation, it achieves very good performance.
- As a **virtualizer**, QEMU achieves near native performances by executing the guest code directly on the host CPU. QEMU supports virtualization when executing under the Xen hypervisor or using the KVM kernel module in Linux. When using KVM, QEMU can virtualize x86, server and embedded PowerPC, and S390 guests.

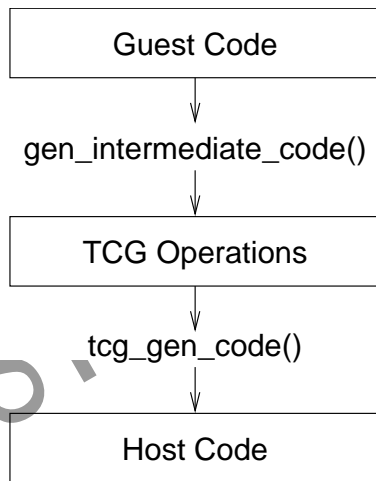


QEMU Internals

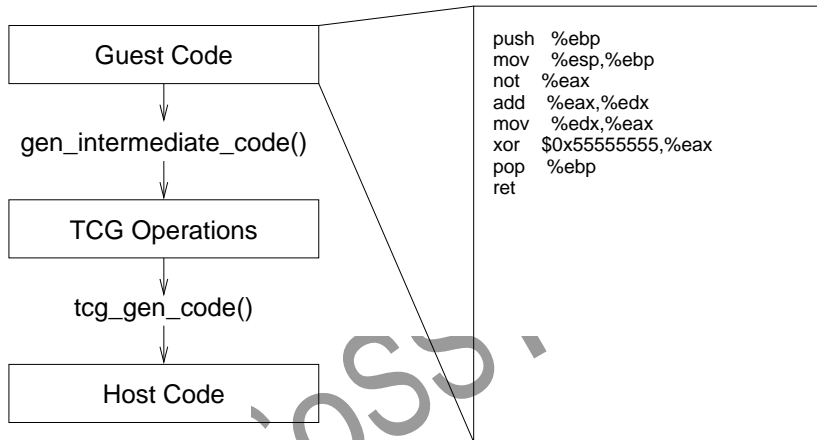


QEMU-Code Translation

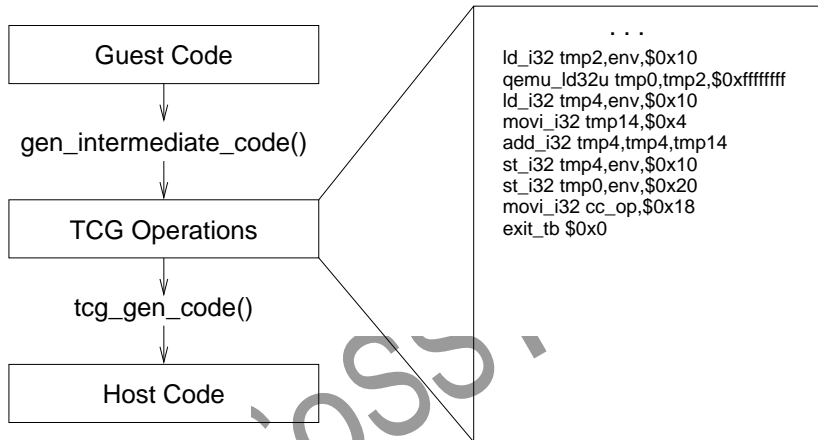
- QEMU uses an intermediate form.
- Frontends are in `target-*/`, includes `alpha`, `arm`, `cris`, `i386`, `m68k`, `mips`, `ppc`, `sparc`, etc.
- Backends are in `tcg/`, includes `arm/`, `hppa/`, `i386/`, `ia64/`, `mips/`, `ppc/`, `ppc64/`, `s390/`, `sparc/`, `tcg.c`, `tcg.h`, `tcg-opc.h`, `tcg-op.h`, `tcg-runtime.h`



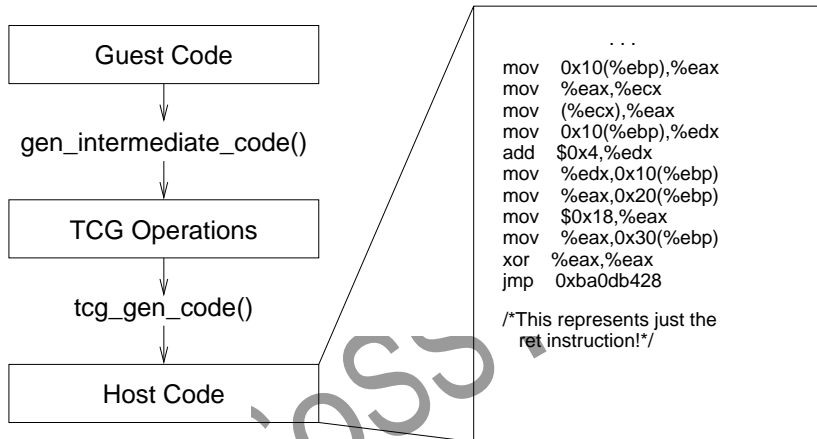
QEMU-Code Translation



QEMU-Code Translation



QEMU-Code Translation



QEMU-Code Base

- TranslationBlock structure in `translate-all.h`
- Translation cache is code gen buffer in `exec.c`
- `cpu-exec()` in `cpu-exec.c` orchestrates translation and block chaining.
- `target-*/translate.c`: guest ISA specific code.
- `tcg-*/*/`: host ISA specific code.
- `linux-user/*`: Linux usermode specific code.
- `vl.c`: Main loop for system emulation.
- `hw/*`: Hardware, including video, audio, and boards.

QEMU use cases

- Malware analysis
- Dynamic binary code instrumentation
- System wide taint analysis
- System wide data lifetime tracking
- Being part of KVM
- Execution replay
- ...

GOSS

Outline

1 Basic Concepts

2 QEMU

3 PIN

4 Summary

GOSS

PIN

- Pin is a tool for the instrumentation of programs. It supports Linux* and Windows* executables for IA-32, Intel(R) 64, and IA-64 architectures.
- Pin allows a tool to insert arbitrary code (written in C or C++) in arbitrary places in the executable. The code is added dynamically while the executable is running. This also makes it possible to attach Pin to an already running process.



Credit: The rest of the slides are compiled from Intel's PIN tutorial

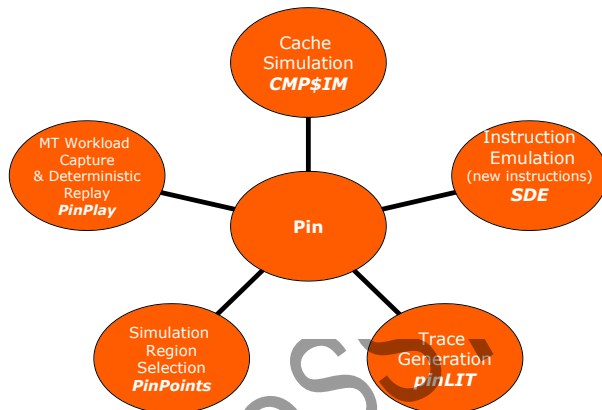
Advantages of Pin Instrumentation

- Easy-to-use Instrumentation:
 - ▶ Uses dynamic instrumentation
 - ★ Does not need source code, recompilation, post-linking
- Programmable Instrumentation:
 - ▶ Provides rich APIs to write in C/C++ your own instrumentation tools (called Pintools)
- Multiplatform:
 - ▶ Supports x86, x86-64, Itanium
 - ▶ Supports Linux, Windows
- Robust:
 - ▶ Instruments real-life applications: Database, web browsers,...
 - ▶ Instruments multithreaded applications
 - ▶ Supports signals
- Efficient:
 - ▶ Applies compiler optimizations on instrumentation code

Pin Instrumentation Capabilities

- ➊ Replace application functions with your own
 - ▶ Call the original application function from within your replacement function
- ➋ Fully examine any application instruction, insert a call to your instrumenting function to be executed whenever that instruction executes
 - ▶ Pass parameters to your instrumenting function from a large set of supported parameters
 - ★ Register values (including EIP), Register values by reference (for modification)
 - ★ Memory address read/written by the instruction
 - ★ Full register context
 - ★ ...
- ➌ Track function calls including syscalls and examine/change arguments
- ➍ Track application threads
- ➎ Intercept signals
- ➏ Instrument a process tree
- ➐ Many other capabilities ...

Example Pin-tools



Using Pin

- Launch and instrument an application:

```
$pin -t pintool.so - - application
```

- 1 instrumentation engine (provided)
- 2 instrumentation tool (write your own, or use a provided sample)

- Attach to a running process, and instrument application:

```
$pin -t pintool.so -pid 1234
```

Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```

GOSS

Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```



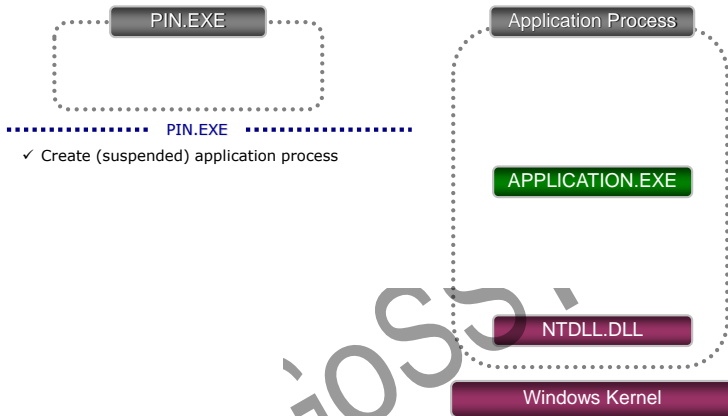
PIN.EXE

The diagram shows a dashed rectangular box representing a process. A small dark gray rounded rectangle with the text 'PIN.EXE' is attached to the top-left corner of the dashed box.

GOSS

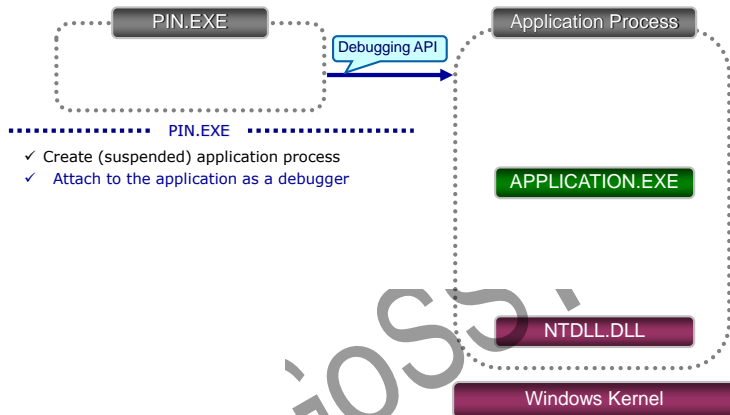
Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```



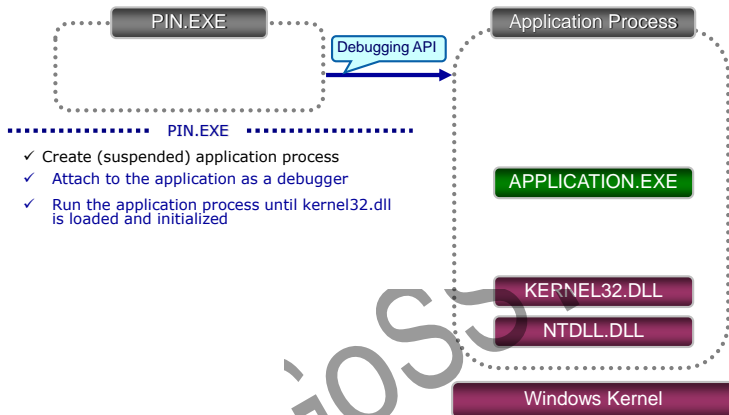
Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```



Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```



Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```

PIN.EXE

Application Process

PIN.EXE

APPLICATION.EXE

KERNEL32.DLL

NTDLL.DLL

Windows Kernel

- ✓ Create (suspended) application process
- ✓ Attach to the application as a debugger
- ✓ Run the application process until kernel32.dll is loaded and initialized
- ✓ Detach from the application process

Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```

PIN.EXE

PIN.EXE

Application Process

Pin Boot Routine

APPLICATION.EXE

KERNEL32.DLL

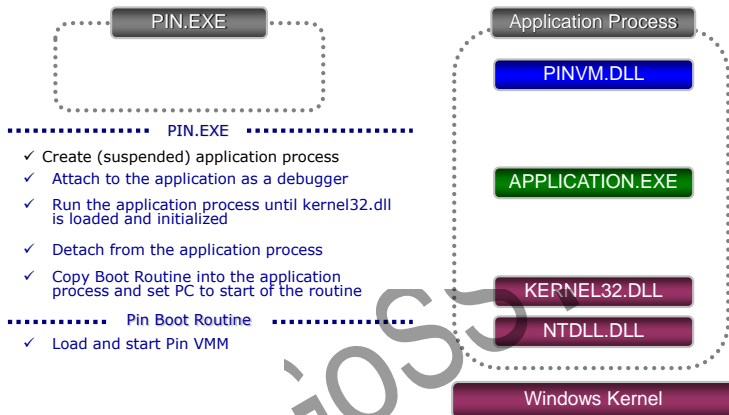
NTDLL.DLL

Windows Kernel

- ✓ Create (suspended) application process
- ✓ Attach to the application as a debugger
- ✓ Run the application process until kernel32.dll is loaded and initialized
- ✓ Detach from the application process
- ✓ Copy Boot Routine into the application process and set PC to start of the routine

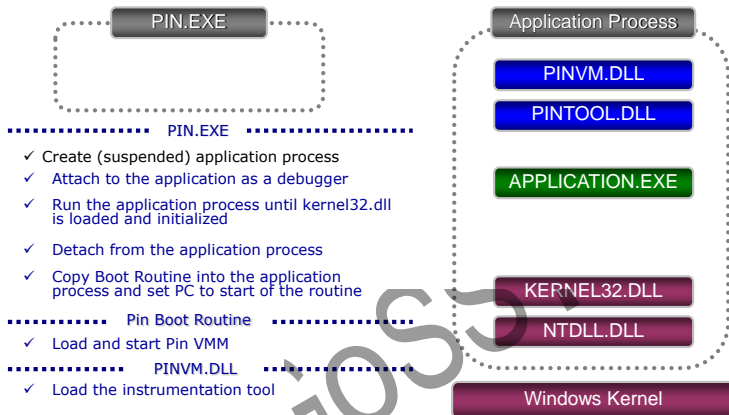
Launch of the Instrumented Process

`pin -t pintool.dll -- application.exe`



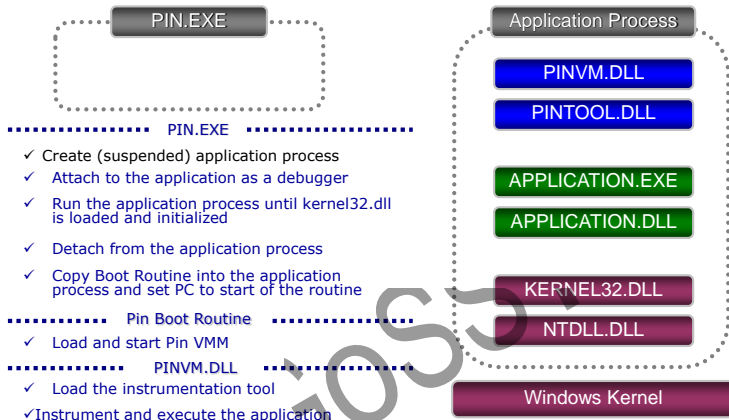
Launch of the Instrumented Process

```
pin -t pintool.dll -- application.exe
```



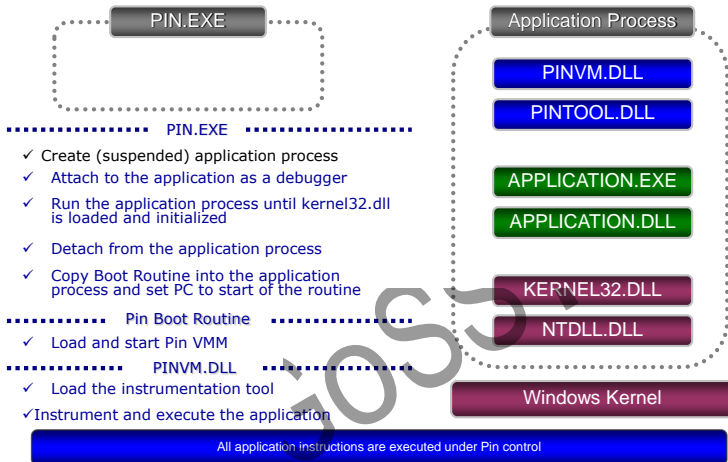
Launch of the Instrumented Process

`pin -t pintool.dll -- application.exe`



Launch of the Instrumented Process

`pin -t pintool.dll -- application.exe`



Pin Instrumentation APIs

- Basic APIs are architecture independent:
 - ▶ Provide common functionalities like determining:
 - ★ Control-flow changes
 - ★ Memory accesses
- Architecture-specific APIs
 - ▶ e.g., Info about opcodes and operands
- Call-based APIs:
 - ▶ Instrumentation routines
 - ▶ Analysis routines

JOSS

Instrumentation vs. Analysis

- Instrumentation routines define where instrumentation is inserted
 - ▶ e.g., before instruction
 - ▶ Occurs first time an instruction is executed
- Analysis routines define what to do when instrumentation is activated
 - ▶ e.g., increment counter
 - ▶ Occurs every time an instruction is executed

goss

ManualExamples/itrace.cpp

```
#include <stdio.h>
#include "pin.h"
FILE * trace;

void printip(void *ip) { fprintf(trace, "%p\n", ip); }

void Instruction(INS ins, void *v) {
    INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)printip, IARG_INST_PTR, IARG_END);
}

void Fini(INT32 code, void *v) { fclose(trace); }

int main(int argc, char * argv[]) {
    trace = fopen("itrace.out", "w");
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```


Examples of Arguments to Analysis Routine

- IARG_INST_PTR
 - ▶ Instruction pointer (program counter) value
- IARG_UINT32 <value>
 - ▶ An integer value
- IARG_REG_VALUE <register name>
 - ▶ Value of the register specified
- IARG_BRANCH_TARGET_ADDR
 - ▶ Target address of the branch instrumented
- IARG_MEMORY_READ_EA
 - ▶ Effective address of a memory read
- And many more ... (refer to the Pin manual for details)

Pintool Example: Instruction trace

- Need to pass the ip argument to the printip analysis routine

```
printip(ip)
sub      $0xff, %edx
printip(ip)
cmp      %esi, %edx
printip(ip)
jle      <L1>
printip(ip)
mov      $0x1, %edi
printip(ip)
add      $0x10, %eax
```

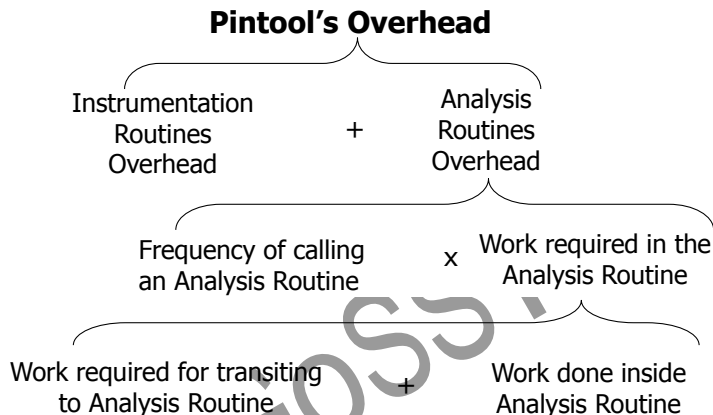
Running itrace tool

```
$ /opt/pin/pin
-t /opt/pin/source/tools/ManualExamples/
obj-intel64/itrace.so
-- /bin/ls

(...)
```

```
$ head itrace.out
0x7f907b188af0
0x7f907b188af3
0x7f907b189120
0x7f907b189121
0x7f907b189124
```

Reducing the Pintool's Overhead



Slower Instruction Counting

```
counter++;  
sub    $0xff, %edx  
counter++;  
cmp    %esi, %edx  
counter++;  
jle    <L1>  
counter++;  
mov    $0x1, %edi  
counter++;  
add    $0x10, %eax
```

Faster Instruction Counting

Counting at BBL level

```
counter += 3
sub    $0xff, %edx
cmp    %esi, %edx
jle    <L1>
```

```
counter += 2
mov    $0x1, %edi

add    $0x10, %eax
```

Counting at Trace level

```
sub    $0xff, %edx
cmp    %esi, %edx
jle    <L1>_____
```

```
mov    $0x1, %edi
add    $0x10, %eax
counter += 5
```

```
counter+=3
```

L1

Writing your own Pintool

- It's easier to modify one of the existing tools, and re-use the existing makefile
- Install PIN package in your home directory, and work from there
 - ▶ `/opt/pin-<version>-<architecture>-<os>.tar.gz`

GOSS

Outline

1 Basic Concepts

2 QEMU

3 PIN

4 Summary

GOSS

Summary

QEMU
open source processor emulator

Valgrind

The logo for Pin Dyn Inst, featuring the word "Pin" in a blue, stylized font inside a blue square, followed by "Dyn" in a large, blue, serif font, and "Inst" in a smaller, blue, serif font.

References

- <http://en.wikipedia.org/wiki/QEMU>
- http://wiki.qemu.org/Main_Page
- <http://valgrind.org/>
- <http://www.pintool.org/>
- <http://www.dyninst.org/>

JOSS